

Homework 4

Shao-Ting Chiu (UIN:433002162)

9/28/22

Description

- Course: STAT638, 2022 Fall

Read Chapter 4 in Hoff.

Then, do the following exercises in Hoff: 4.1, 4.2, 4.6, 4.8

All datasets in the Hoff book can be downloaded from <https://pdhoff.github.io/book/> (Links to an external site.).

- Deadline: Oct 4 by 12:01pm
-

Computational Environment Setup

Third-party libraries

```
%matplotlib inline
import sys # system information
import matplotlib # plotting
import scipy # scientific computing
import random
import pandas as pd # data managing
from scipy.special import comb
from scipy import stats as st
from scipy.special import gamma
import numpy as np
import matplotlib.pyplot as plt
# Matplotlib setting
```

```
plt.rcParams['text.useTex'] = True
matplotlib.rcParams['figure.dpi']= 300
np.random.seed(20220928) # Consistent random effect
```

Version

```
print(sys.version)
print(matplotlib.__version__)
print(scipy.__version__)
print(np.__version__)
print(pd.__version__)
```

```
3.8.12 (default, Oct 22 2021, 18:39:35)
[Clang 13.0.0 (clang-1300.0.29.3)]
3.3.1
1.5.2
1.19.1
1.1.1
```

Problem 4.1

Posterior comparisons: Reconsider the sample survey in Exercise 3.1. Suppose you are interested in comparing the rate of support in that county to the rate in another county. Suppose that a survey of sample size 50 was done in the second county, and the total number of people in the sample who supported the policy was 30. Identify the posterior distribution of θ_2 assuming a uniform prior. Sample 5000 values of each of θ_1 and θ_2 from their posterior distributions and estimate $Pr(\theta_1 < \theta_2 | \text{the data and prior})$.

- Prior:

$$- \theta \sim \text{beta}(1, 1)$$

- Model:

$$- p(\sum Y = n | \theta) = \binom{N}{n} \theta^n (1 - \theta)^{N-n}$$

- Posterior distribution (Hoff 2009, 580:37):

$$\begin{aligned}
& - \theta | \sum_{i=1}^N Y_i = n \sim \text{Beta}(\text{beta}(1 + n, 1 + N - n)) \\
& * \theta_1 \sim \text{Beta}(1 + 57, 1 + 100 - 57) = \text{Beta}(58, 44) \\
& * \theta_2 \sim \text{Beta}(1 + 30, 1 + 50 - 30) = \text{Beta}(31, 21)
\end{aligned}$$

```

N = 5000
t1s = st.beta.rvs(58, 44, size=N)
t2s = st.beta.rvs(31, 21, size=N)
p_t2bigger = np.mean(t1s < t2s)

# Display
pd.DataFrame({"Item": ["Pr(theta_1 < theta_2 | the data and prior)"],\
                  "Value": [p_t2bigger]})

```

	Item	Value
0	Pr(theta_1 < theta_2 the data and prior)	0.6292

Problem 4.2

Tumor count comparisons: Reconsider the tumor count data in [Exercise 3.3](#):

(a)

For the prior distribution given in part (a) of that exercise, obtain $Pr(\theta_B < \theta_A | y_A, y_B)$ via Monte Carlo sampling.

- Prior distribution

$$\begin{aligned}
& - \theta_A | y_A \sim \text{gamma}(120 + 117, 10 + 10) = \text{gamma}(237, 20) \\
& - \theta_B | y_B \sim \text{gamma}(12 + 113, 1 + 13) = \text{gamma}(125, 14)
\end{aligned}$$

```

N = 5000
theta_A = st.gamma.rvs(237, scale=1/20, size=N)
theta_B = st.gamma.rvs(125, scale=1/14, size=N)
res = np.mean(theta_B < theta_A)

# Display
pd.DataFrame({"Item": ["Pr(theta_B < theta_A | y_A, y_B)"],\
                  "Value": [res]})

```

Item	Value
0 Pr($\theta_B < \theta_A \mid y_A, y_B$)	0.995

(b)

For a range of values of n_0 , obtain $Pr(\theta_B < \theta_A \mid y_A, y_B)$ for $\theta_A \sim \text{gamma}(120, 10)$ and $\theta_B \sim \text{gamma}(12 \times n_0, n_0)$. Describe how sensitive the conclusions about the event $\{\theta_B < \theta_A\}$ are to the prior distribution on θ_B .

In Figure 1, the n_0 decreases the probability of $p(\{\theta_B < \theta_A\})$ with linear effect. From small to large n_0 , the prior distribution of θ_B keep influences the result.

```
def get_A_bigger(n0, N=5000):
    theta_A = st.gamma.rvs(120 + 117, scale=1/(10+10), size=N)
    theta_B = st.gamma.rvs(12*n0+113, scale=1/(n0+13), size=N)
    res = np.mean(theta_B < theta_A)
    return res

n0s = np.arange(1,50, 1)
ress = [get_A_bigger(n0) for n0 in n0s]

# Plotting
fig, ax = plt.subplots()
ax.plot(n0s, ress, "o", color="k")
ax.set_xlabel("$n_0$")
ax.set_ylabel("$p(\theta_B < \theta_A \mid y_A, y_B, n_0)$");
```




Figure 1: The effect of n_0

(c)

Repeat part (a) and (b), replacing the event $\{\theta_B < \theta_A\}$ with the event $\{\tilde{Y}_B < \tilde{Y}_A\}$, where \tilde{Y}_A and \tilde{Y}_B are samples from the posterior predictive distribution.

Part I (a)

- $\tilde{Y}_A \sim \text{nbinom}(a + \sum Y^{(A)}, b + n) = \text{nbinom}(120 + 117, 10 + 10) = \text{nbinom}(237, 20)$
- $\tilde{Y}_B \sim \text{nbinom}(a + \sum Y^{(B)}, b + n) = \text{nbinom}(12 + 113, 1 + 13) = \text{nbinom}(125, 14)$
- Find $p(\tilde{Y}_B < \tilde{Y}_A | y_A, y_B)$
- Use `scipy.stats.nbinom`¹

- $p = 1 - \frac{1}{b+n+1}$
- $n = a + \sum Y$
- $k = \tilde{Y}$

¹<https://docs.scipy.org/doc/scipy/reference/generated/scipy.stats.nbinom.html#scipy.stats.nbinom>

```

N = 5000
tilde_y_A = st.nbinom.rvs(120+117, 1 - 1/(10+10+1), size=N)
tilde_y_B = st.nbinom.rvs(12+113, 1 - 1/(1+13+1), size=N)
p2 = np.mean(tilde_y_B < tilde_y_A)

# Display
# Display
pd.DataFrame({"Item": ["Pr(tildeY_B < tildeY_A | y_A, y_B)"],\
                 "Value": [p2]})

```

Item	Value
0 Pr(tildeY_B < tildeY_A y_A, y_B)	0.6934

Part II (b)

In Figure 2, the n_0 has nonlinear negative effect on the probability $p(\{\tilde{Y}_B < \tilde{Y}_A\})$ with the prior formation of n_0 .

```

def get_A_bigger(n0, N=5000):
    tilde_y_A = st.nbinom.rvs(120+117, 1 - 1/(10+10+1), size=N)
    tilde_y_B = st.nbinom.rvs(12*n0+113, 1 - 1/(n0+13+1), size=N)
    res = np.mean(tilde_y_B < tilde_y_A)
    return res

n0s = np.arange(1,50, 1)
ress = [get_A_bigger(n0) for n0 in n0s]

# Plotting
fig, ax = plt.subplots()
ax.plot(n0s, ress, "o", color="k")
ax.set_xlabel("$n_0$")
ax.set_ylabel("$p(\tilde{Y}_B < \tilde{Y}_A | y_A, y_B, n_0)$");

```




Figure 2: The effect of n_0

Problem 4.6

Non-informative prior distributions: Suppose for a binary sampling problem we plan on using a uniform, or $beta(1, 1)$, prior for the population proportion θ . Perhaps our reasoning is that this represents “no prior information about θ .” However, some people like to look at proportions on the log-odds scale, that is, they are interested in $\gamma = \log \frac{\theta}{1-\theta}$. Via Monte Carlo sampling or otherwise, find the prior distribution for γ that is induced by the uniform prior for θ . Is the prior informative about γ ?

Part I: Analytical Approach

- $\gamma = g(\theta) = \log \frac{\theta}{1-\theta}$
- $\theta = g^{-1}(\gamma) = \frac{e^\gamma}{1+e^\gamma}$

$$p_\gamma(\gamma) = \underbrace{p_\theta(g^{-1}(\gamma))}_{=1: \theta \sim \text{uniform}(0,1)} \times \left| \frac{dg^{-1}(\gamma)}{d\gamma} \right| \quad (1)$$

$$= \left| \frac{e^\gamma}{(1+e^\gamma)^2} \right| \quad (2)$$

Part II: Monte Carlo Approach


```
def true_gamma_pdf(y):
    return np.absolute(np.exp(y)/(1+np.exp(y))**2)

ths = st.beta.rvs(1,1,size=1000)
gammas = np.linspace(-10,10,100)

ys_true = [true_gamma_pdf(g) for g in gammas]
ys = [np.log(th/(1-th)) for th in ths]

fig, ax = plt.subplots();
ax.hist(ys, weights=np.ones_like(ys)/len(ys));
ax.plot(gammas, ys_true, label="Analytical solution");
ax.set_xlabel("$\\gamma = \\log \\frac{\\theta}{1-\\theta}$");
ax.set_ylabel("$p(\\gamma)$");
ax.legend()
```

<matplotlib.legend.Legend at 0x1339db790>



Problem 4.8

More posterior predictive checks: Let θ_A and θ_B be the average number of children of men in their 30s with and without bachelor's degrees, respectively.

(a)

Using a Poisson sampling model, a $\text{gamma}(2, 1)$ prior for each θ and the data in the files `menchild30bach.dat` and `menchild30nobach.dat`, obtain 5000 samples of \bar{Y}_A and \bar{Y}_B from the posterior predictive distribution of the two samples. Plot the Monte Carlo approximations to these two posterior predictive distributions. (data available in [Appendix](#))

```
# Data
data_A = np.loadtxt("data/menchild30bach.dat")
data_B = np.loadtxt("data/menchild30nobach.dat")

# Statistics
suma = np.sum(data_A)
sumb = np.sum(data_B)
Na = len(data_A)
Nb = len(data_B)

# Display
pd.DataFrame({
    "Properties": ["Sum", "N (number of samples)"],
    "A": [suma, Na],
    "B": [sumb, Nb]
})
```

	Properties	A	B
0	Sum	54.0	305.0
1	N (number of samples)	58.0	218.0

The predictive distribution is

- $\bar{Y}_A | y_A \sim \text{nbinom}(2 + 54, 1 + 58)$
- $\bar{Y}_B | y_B \sim \text{nbinom}(2 + 305, 1 + 218)$

```
N = 5000
bar_ya = st.nbinom.rvs(2+54, 1 - 1/(1+58+1), size=N)
```


```

bar_yb = st.nbinom.rvs(2+305, 1 - 1/(1+218+1), size=N)

# Display
fig, ax = plt.subplots();
[ax.hist(ys, weights=np.ones_like(ys)/len(ys), label=n, alpha=0.7, bins=np.arange(0,8,1))
ax.set_xlabel("$N$")
ax.set_ylabel("Probability")
ax.legend()

```

<matplotlib.legend.Legend at 0x1304210d0>



(b)

Find 95% quantile-based posterior confidence intervals for $\theta_B - \theta_A$ and $\tilde{Y}_B - \tilde{Y}_A$. Describe in words the differences between the two populations using these quantities and the plots in (a), along with any other results that may of interest to you.

(c)

Obtain the empirical distribution of the data in group B . Compare this to the Poisson distribution with mean $\hat{\theta} = 1.4$. Do you think the Poisson model is a good fit? Why or Why not?

(d)

For each of the 5000 θ_B -values you sampled, sample $n_B = 218$ Poisson random variables and count the number of 0s and the number of 1s in each of the 5000 simulated datasets. You should now have two sequences of length 5000 each, one sequence counting the number of people having zero children for each of the 5000 posterior predictive datasets, the other counting the number of people with one child. Plot the two sequences against one another (one on the x -axis, one on the y -axis). Add to the plot a point marking how many people in the observed dataset had zero children and one child. Using this plot, describe the adequacy of the Poisson model.

Appendix

Data set in [Problem 4.8](#)

```
print("menchild30bach.dat:\n", data_A)
print("menchild30nobach.dat:\n", data_B)
```

menchild30bach.dat:

```
[1. 0. 0. 1. 2. 2. 1. 5. 2. 0. 0. 0. 0. 0. 0. 1. 1. 1. 0. 0. 0. 1. 1. 2.
 1. 3. 2. 0. 0. 3. 0. 0. 0. 2. 1. 0. 2. 1. 0. 0. 1. 3. 0. 1. 1. 0. 2. 0.
 0. 2. 2. 1. 3. 0. 0. 0. 1. 1.]
```

menchild30nobach.dat:

```
[2. 2. 1. 1. 2. 2. 1. 2. 1. 0. 2. 1. 1. 2. 0. 2. 2. 0. 2. 1. 0. 0. 3. 6.
 1. 6. 4. 0. 3. 2. 0. 1. 0. 0. 0. 3. 0. 0. 0. 0. 0. 1. 0. 4. 2. 1. 0. 0.
 1. 0. 3. 2. 5. 0. 1. 1. 2. 1. 2. 1. 2. 0. 0. 0. 2. 1. 0. 2. 0. 2. 4. 1.
 1. 1. 2. 0. 1. 1. 1. 1. 0. 2. 3. 2. 0. 2. 1. 3. 1. 3. 2. 2. 3. 2. 0. 0.
 0. 1. 0. 0. 0. 1. 2. 0. 3. 3. 0. 1. 2. 2. 2. 0. 6. 0. 0. 0. 2. 0. 1. 1.
 1. 3. 3. 2. 1. 1. 0. 1. 0. 0. 2. 0. 2. 0. 1. 0. 2. 0. 0. 2. 2. 4. 1. 2.
 3. 2. 0. 0. 0. 1. 0. 0. 1. 5. 2. 1. 3. 2. 0. 2. 1. 1. 3. 0. 5. 0. 0. 2.
 4. 3. 4. 0. 0. 0. 0. 0. 0. 2. 2. 0. 0. 2. 0. 0. 1. 1. 0. 2. 1. 3. 3. 2.
 2. 0. 0. 2. 3. 2. 4. 3. 3. 4. 0. 3. 0. 1. 0. 1. 2. 3. 4. 1. 2. 6. 2. 1.
 2. 2.]
```

References

Hoff, Peter D. 2009. *A First Course in Bayesian Statistical Methods*. Vol. 580. Springer.