

Homework 6

Shao-Ting Chiu (UIN:433002162)

10/14/22

Table of contents

Description	1
Computational Enviromnent Setup	2
Third-party libraries	2
Version	2
Problem 6.1	3
(a)	3
(b)	3
(c)	4
(d)	5
Problem External	7
(a)	7
(b)	8
(c)	9
(d)	11

Description

- Course: STAT638, 2022 Fall

Read Hoff (2009, ch. 6). Then, do Hoff (2009, Exercise 6.1). You may assume that θ and γ are a priori independent, and that Y_A and Y_B are conditionally independent given θ and γ .

Computational Environment Setup

Third-party libraries

```
%matplotlib inline
import sys # system information
import matplotlib # plotting
import scipy # scientific computing
import random
import pandas as pd # data managing
from scipy.special import comb
from scipy import stats as st
from scipy.special import gamma
import numpy as np
import matplotlib.pyplot as plt
from itertools import permutations
# Matplotlib setting
plt.rcParams['text.usetex'] = True
matplotlib.rcParams['figure.dpi']= 300
np.random.seed(20220928) # Consistent random effect
```

Version

```
print(sys.version)
print(matplotlib.__version__)
print(scipy.__version__)
print(np.__version__)
print(pd.__version__)
```

3.8.14 (default, Sep 6 2022, 23:26:50)

[Clang 13.1.6 (clang-1316.0.21.2.5)]

3.3.1

1.5.2

1.19.1

1.1.1

Problem 6.1

Poisson population comparisons: Let's reconsider the number of children data of Exercise 4.8. We'll assume Poisson sampling models for the two groups as before, but now we'll parameterize θ_A and θ_B as $\theta_A = \theta, \theta_B = \theta \times \gamma$. In the parameterization, γ represents the relative rate $\frac{\theta_B}{\theta_A}$. Let $\theta \sim \text{gamma}(a_\theta, b_\theta)$ and let $\gamma \sim \text{gamma}(a_\gamma, b_\gamma)$.

(a)

Are θ_A and θ_B independent or dependent under this prior distribution? In what situations is such a joint prior distribution justified?

$$\text{Cov}(\theta_A, \theta_B) = E[\theta_A \theta_B] - E[\theta_A]E[\theta_B] \quad (1)$$

$$= E[\theta \theta \gamma] - E[\theta]E[\theta \gamma] \quad (2)$$

$$= E[\theta^2 \gamma] - E[\theta]E[\theta \gamma] \quad (3)$$

$$= E[\theta^2]E[\gamma] - E[\theta]^2 E[\gamma] \quad (4)$$

$$= E[\gamma](E[\theta^2] - E[\theta]^2) \quad (5)$$

$$= E[\gamma]\text{Var}[\theta] \quad (6)$$

$$= \frac{a_\gamma}{b_\gamma} \frac{a_\theta}{b_\theta^2} \quad (7)$$

$$\neq 0 \quad (8)$$

Because $\text{Cov}(\theta_A, \theta_B) > 0$, θ_A and θ_B are dependent.

(b)

Obtain the form of the full conditional distribution of θ given y_A, y_B and γ .

$$p(\theta|y_A, y_B, \gamma) = \frac{p(y_A, y_B|\theta, \gamma)p(\theta|\gamma)}{p(y_A, y_B|\gamma)} \quad (9)$$

$$= \frac{p(y_A|\theta, \gamma)p(y_B|\theta, \gamma)p(\theta|\gamma)}{p(y_A|\gamma)p(y_B|\gamma)} \quad (10)$$

$$= \frac{p(y_A|\theta)p(y_B|\theta, \gamma)p(\theta)}{p(y_A)p(y_B|\gamma)} \quad (11)$$

$$\propto p(y_A|\theta)p(y_B|\theta, \gamma)p(\theta) \quad (12)$$

$$= \prod_{i=1}^{n_A} \text{poisson}(y_{Ai}; \theta) \times \prod_{i=1}^{n_B} \text{poisson}(y_{Bi}; \theta\gamma) \times \text{gamma}(\theta; a_\theta, b_\theta) \quad (13)$$

$$= \theta^{\sum_{i=1}^{n_A} y_{Ai}} e^{-n_A \theta} \times (\theta\gamma)^{\sum_{i=1}^{n_B} y_{Bi}} e^{-n_B \theta\gamma} \times \theta^{a_\theta-1} e^{-b_\theta \theta} \quad (14)$$

$$= \theta^{n_A \bar{y}_A} e^{-n_A \theta} \times (\theta\gamma)^{n_B \bar{y}_B} e^{-n_B \theta\gamma} \times \theta^{a_\theta-1} e^{-b_\theta \theta} \quad (15)$$

$$= \theta^{n_A \bar{y}_A} e^{-n_A \theta} \times (\theta)^{n_B \bar{y}_B} e^{-n_B \theta\gamma} \times \theta^{a_\theta-1} e^{-b_\theta \theta} \quad (16)$$

$$\propto \theta^{n_A \bar{y}_A + n_B \bar{y}_B + a_\theta - 1} e^{-(n_A + n_B \gamma + b_\theta) \theta} \quad (17)$$

Therefore, $p(\theta|y_A, y_B, \gamma) \sim \text{gamma}(\theta|n_A \bar{y}_A + n_B \bar{y}_B + a_\theta, n_A + n_B \gamma + b_\theta)$

(c)

Obtain the form of the full conditional distribution of γ given y_A, y_B and θ .

$$p(\gamma|y_A, y_B, \theta) = \frac{p(y_A, y_B|\gamma, \theta)p(\gamma|\theta)}{p(y_A, y_B|\theta)} \quad (18)$$

$$\propto p(y_A, y_B|\gamma, \theta)p(\gamma|\theta) \quad (19)$$

$$= p(y_A|\gamma, \theta)p(y_B|\gamma, \theta)p(\gamma) \quad (20)$$

$$= p(y_A|\theta)p(y_B|\gamma, \theta)p(\gamma) \quad (21)$$

$$= \prod_{i=1}^{n_A} p(y_{Ai}|\theta) \prod_{i=1}^{n_B} p(y_{Bi}|\gamma, \theta)p(\gamma) \quad (22)$$

$$= \theta^{n_A \bar{y}_A} e^{-n_A \theta} (\gamma\theta)^{n_B \bar{y}_B} e^{-n_B \gamma\theta} \gamma^{a_\gamma-1} e^{-b_\gamma \gamma} \quad (23)$$

$$\propto \gamma^{n_B \bar{y}_B + a_\gamma - 1} e^{-n_A \theta - n_B \gamma\theta - b_\gamma \gamma} \quad (24)$$

$$\propto \gamma^{n_B \bar{y}_B + a_\gamma - 1} e^{-(n_B \theta + b_\gamma) \gamma} \quad (25)$$

Therefore, $p(\gamma|y_A, y_B, \theta) \sim \text{Gamma}(n_B \bar{y}_B + a_\gamma, n_B \theta + b_\gamma)$

(d)

Set $a_\theta = 2$ and $b_\theta = 1$. Let $a_\gamma = b_\gamma \in \{8, 16, 32, 64, 128\}$. For each of these five values, run a Gibbs sampler of at least 5000 iterations and obtain $E[\theta_B - \theta_A | \mathcal{Y}_A, \mathcal{Y}_B]$. Describe the effects of the prior distribution for γ on the results.

```
class GibbSampler:
    def __init__(self, a_t, b_t, a_g, b_g):
        # Theta
        self.a_t = 2
        self.b_t = 1
        # Gamma
        self.a_g = a_g
        self.b_g = b_g
        self.theta = st.gamma(a_t, scale=1/b_t)
        self.gamma = st.gamma(a_g, scale=1/b_g)

    def expection(self, obs_func, dataA, dataB, n_sampling):
        # Sample number
        nA=len(dataA)
        nB= len(dataB)
        # Sample mean
        mA = np.mean(dataA)
        mB = np.mean(dataB)

        thetaAs = np.zeros(n_sampling)
        thetaBs = np.zeros(n_sampling)

        # initial
        theta_pre = st.gamma.rvs(self.a_t, scale=(self.b_t)**-1)
        gamma_pre = st.gamma.rvs(self.a_g, scale=(self.b_g)**-1)

        # Gibbs iteration
        thetaAs[0] = theta_pre
        thetaBs[0] = theta_pre*gamma_pre

        for i in range(1,n_sampling):
            # posterior
            theta = st.gamma.rvs(nA*mA + nB*mB + self.a_t, scale= (nA + nB*gamma_pre + self.a_t)**-1)
            theta_pre = theta
            gamma = st.gamma.rvs(nB*mB + self.a_g, scale= (nB*theta_pre + self.b_g)**-1)
```

```

        gamma_pre = gamma

        # transformation
        thetaAs[i] = theta
        thetaBs[i] = theta*gamma

    # Apply obs
    obs = [obs_func(thetaAs[i], thetaBs[i]) for i in range(0,n_sampling)]

    return np.mean(obs)

def obs_func(theta_A, theta_B):
    return theta_B - theta_A

vals = np.array([8,16,32,64,128], dtype=float)

priors = [{\
    "a_t": 2.0,
    "b_t": 1.0,
    "a_g": a,
    "b_g": a
} for a in vals]

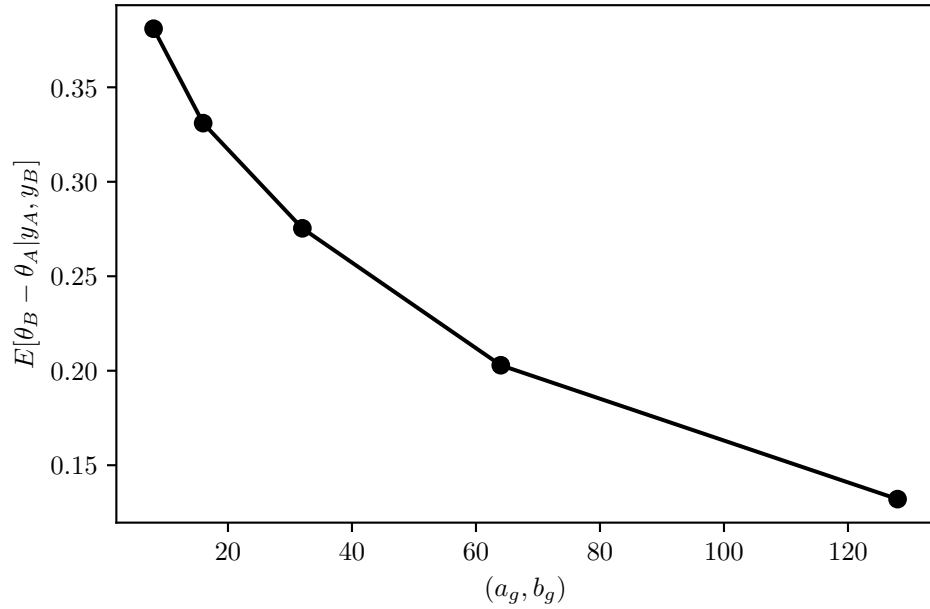
obs = {
    "obs_func": obs_func,
    "dataA": np.loadtxt("data/menchild30bach.dat"),
    "dataB": np.loadtxt("data/menchild30nobach.dat"),
    "n_sampling": 5000
}

exps = np.zeros(len(vals))

for i in range(0, len(vals)):
    g = GibbsSampler(**priors[i])
    exps[i] = g.expectation(**obs)

plt.plot(vals, exps, "-o",color="black");
plt.xlabel("$a_{g}, b_{g}$");
plt.ylabel("$E[\theta_{B}-\theta_{A} \mid y_A, y_B]$");

```



Problem External

Also complete the following problem: We would like to study the survival times after patients receive a new cancer treatment. We observe the following survival times (in years) for 6 patients: 3, 5, x , 4, x , x . Here, x denotes a censored observation, meaning that the respective patient survived for more than 5 years after the treatment (which is when the study ended). We consider the following model:

$$Y_i = \begin{cases} Z_i, & Z_i \leq c \\ \times, & Z_i > c \end{cases}, i = 1, \dots, n \quad (26)$$

$$Z_1, \dots, Z_n | \theta \sim^{iid} \text{Exponential}(\theta)$$

$$\theta \sim \text{Gamma}(a, b)$$

We have $a = 1$, $b = 4$, $c = 5$, and $n = 6$.

(a)

Find the full-conditional distribution (FCD) of θ

$$p(\theta|y, z) \propto \underbrace{p(y|\theta, z)}_{=1} p(\theta|z) \quad (27)$$

$$= p(\theta|z) \quad (28)$$

$$= \frac{p(z|\theta)p(\theta)}{p(z)} \quad (29)$$

$$\propto p(z|\theta)p(\theta) \quad (30)$$

$$= \prod_{i=1}^{n=6} p(z_i|\theta)p(\theta) \quad (31)$$

- $p(z_i|\theta) \propto \theta e^{-\theta z_i}$
- $p(\theta) \propto \theta^{a-1} e^{-b\theta}$

$$p(\theta|y, z) = \left(\prod_{i=1}^{n=6} \theta e^{-\theta z_i} \right) \theta^{a-1} e^{-b\theta} \quad (32)$$

$$= \theta^n e^{-\theta n \bar{z}} \theta^{a-1} e^{-b\theta} \quad (33)$$

$$= \theta^{n+a-1} e^{-b\theta - \theta n \bar{z}} \quad (34)$$

$$= \theta^{(n+a)-1} e^{-(b+n\bar{z})\theta} \quad (35)$$

Therefore, $\theta \sim \text{gamma}(n + a, b + n\bar{z})$ - where $\bar{z} = \frac{1}{n} \sum_{i=1}^n z_i$

(b)

Find the FCD of each Z_i .

(Hint: For uncensored i , this distribution will be a degenerate point mass; for censored i , the resulting distribution will be a so-called truncated exponential distribution, which is proportional to a exponential density but constrained to lie in an interval. Each FCD does not depend on other Z 's)

$$p(z_i|\theta, y_i = x) = \theta e^{-\theta(z_i - c)} \quad (36)$$

- $p(z_i|\theta, y_i \text{ is uncensored}) = 1$

$$p(z_i|\theta, y_i = x) = p(z_i = c + t|\theta, z_i > c) \quad (37)$$

$$= \frac{p(z_i = c + t \cap z_i > c|\theta)}{p(z_i > c|\theta)} \quad (38)$$

$$= \frac{p(z_i = c + t|\theta)}{p(z_i > c|\theta)} \quad (39)$$

$$= \frac{\theta e^{-\theta(c+t)}}{\int_c^\infty \theta e^{-\theta x} dx} \quad (40)$$

$$= \theta e^{-\theta t} \quad (41)$$

$$= \theta e^{-\theta(z_i - c)} \quad (42)$$

This is the memoryless property of Exponential distribution¹.

In summary,

$$p(z_i|\theta, y_i) = \begin{cases} 1, & y = z_i \\ \theta e^{-\theta(z_i - c)} & y = x \end{cases} \quad (43)$$

(c)

Implement a Gibbs sampler that approximate the joint posterior of θ and Z_1, \dots, Z_n . (For example, you can use `truncdist::rtrunc(3, spec=exp, a=c, rate=theta)` (Use `stats.truncexpon`) to sample from a truncated exponential in R.) Run the sampler for enough iterations such that each of the effective sample sizes for θ and for the three censored Z_i are all greater than 1000. Provide the corresponding trace plots and discuss the mixing of the Markov chain.

- Both sampling converges to a range of stochastic state, that means the sampling is not stuck in certain region. The dependence between two subsequence steps is low.

```
data = np.array([3,5,4], dtype=float)
n = 6
n_cs = n - len(data)
a = 1
b = 4
c = 5
nSamp = 1000
```

¹<https://byjus.com/maths/exponential-distribution/>

```

thetas = np.zeros(nSamp)
zs = np.zeros((nSamp, n_cs))

thetas[0] = 1
zs[0] = [7,10,8]


for i in range(1, nSamp):
    df = np.concatenate([data, zs[i-1]])
    thetas[i] = st.gamma.rvs(n+a, scale= (b+np.sum(df))**-1)
    zs[i] = st.expon.rvs(loc = c, scale= thetas[i], size =3)

# Plotting
fig, ax = plt.subplots(2,1)

ax[0].plot(thetas, "-", color="k")
ax[1].plot(zs[:,0], "-", color="k")

ax[1].set_xlabel("Iteration")
ax[0].set_ylabel("$\\theta$")
ax[1].set_ylabel("Z");

```




- Truncated exponential distribution is sampled by [scipy.stats.truncexpon](#) ²

² $f(x, b) = \frac{\exp(-x)}{1 - \exp(-b)}$. $y = (x - loc)/scale$

(d)

Obtain an approximate 96% posterior credible interval for θ based on the samples from (c).

```
plt.hist(thetas);
```



```
invs = st.mstats.mquantiles(thetas, prob=[0.02, 0.98])
```

```
pd.DataFrame({"96% Credible Interval of theta": ["Left bound", "Right bound"], "Value": invs})
```

96% Credible Interval of theta		Value
0	Left bound	0.085378
1	Right bound	0.444496

Hoff, Peter D. 2009. *A First Course in Bayesian Statistical Methods*. Vol. 580. Springer.