

Homework 8

Shao-Ting Chiu (UIN:433002162)

11/8/22

Table of contents

Description	1
Computational Enviromnent Setup	2
Third-party libraries	2
Version	2
Problem 8.1	3
(a)	3
(b)	3
(c)	4
(d)	5
Problem 8.3	6
(a)	6
(b)	10
(c)	11
(d)	12
(e)	13

Description

- Course: STAT638, 2022 Fall

Read Chapter 8 in the Hoff book. Then do the following exercises in Hoff: 8.1 and 8.3.

Please note some typos in 8.1: All θ_i 's should be θ_j 's.

For 8.1(c), you may find [the law of total \(co-\)variance](#) useful. In addition, remember that all of these laws also hold for conditional distributions (e.g., when conditioning on additional quantities such as μ and τ^2 in all terms on the left- and right-hand side of the equation).

Computational Environment Setup¹

Third-party libraries

```
using Pkg
Pkg.activate("hw8")
using Distributions
using DataFrames
using Turing
using Plots
using DelimitedFiles
using LinearAlgebra
using Statistics
using Turing
using StatsBase
using StatsPlots
import Random
Random.seed!(2022)
```

Activating project at `~/Documents/GitHub/STAT638_Applied-Bayes-Methods/hw/hw8`

Random.TaskLocalRNG()

Version

```
Pkg.status()
VERSION
```

```
Status `~/Documents/GitHub/STAT638_Applied-Bayes-Methods/hw/hw8/Project.toml`
[a93c6f00] DataFrames v1.4.2
[31c24e10] Distributions v0.25.76
[be115224] MCMCDiagnosticTools v0.1.4
[91a5bcdd] Plots v1.35.5
[2913bbd2] StatsBase v0.33.21
[f3b207a7] StatsPlots v0.15.4
[fce5fe82] Turing v0.21.12
[8bb1440f] DelimitedFiles
```

¹I use special character in Julia code. Unfortunately, those are not displayed in PDF version.

v"1.8.2"

Problem 8.1

Components of variance: Consider the hierarchical model where

$$\theta_1, \dots, \theta_m | \mu, \tau^2 \sim i.i.d.\text{normal}(\mu, \tau^2)$$

$$y_{1,j}, \dots, y_{n_j,j} | \theta_j, \sigma^2 \sim i.i.d.\text{normal}(\theta_j, \sigma^2)$$

For this problem, we will eventually compute the following:

- $Var[y_{i,j} | \theta_i, \sigma^2], Var[\bar{y}_{.,j} | \theta_i, \sigma^2], Cov[y_{i_1,j}, y_{i_2,j} | \theta_j, \sigma^2]$
- $Var[y_{i,j} | \mu, \tau^2], Var[\bar{y}_{.,j} | \mu, \tau^2], Cov[y_{i_1,j}, y_{i_2,j} | \mu, \tau^2]$ First, let's use our intuition to guess at the answers:

(a)

Which do you think is bigger, $Var[y_{i,j} | \theta_i, \sigma^2]$ or $Var[y_{i,j} | \mu, \tau^2]$? To guide your intuition, you can interpret the first as the variability of the Y 's when sampling from a fixed group, and the second as the variability in first sampling a group, then sampling a unit from within the group.

- $Var[y_{i,j} | \mu, \tau^2]$ because θ_j is uncertain and the between-group variability create additional uncertainty.

(b)

Do you think $Cov[y_{i_1,j}, y_{i_2,j} | \theta_j, \sigma^2]$ is negative, positive, or zero? Answer the same for $Cov[y_{i_1,j}, y_{i_2,j} | \mu, \tau^2]$. You may want to think about what $y_{i_2,j}$ tells you about $y_{i_1,j}$ if θ_j is known, and what it tells you when θ_j is unknown.

$$Cov[y_{i_1,j}, y_{i_2,j} | \theta_j, \sigma^2]$$

Because $y_{i_1,j}$ and $y_{i_2,j}$ is i.i.d. sampled, I expect $Cov[y_{i_1,j}, y_{i_2,j} | \theta_j, \sigma^2]$ to be zero.

$$Cov[y_{i_1,j}, y_{i_2,j} | \mu, \tau^2]$$

$y_{1,j}$ does tell information about $y_{2,j}$. The covariance $Cov[y_{i_1,j}, y_{i_2,j} | \mu, \tau^2]$ is likely to be positive because values from same θ_j tend to be close together.

(c)

Now compute each of the six quantities above and compare to your answers in (a) and (b). ²

$$Var[y_{i,j} | \theta_i, \sigma^2] = \sigma^2 \quad (1)$$

$$Var[\bar{y}_{\cdot,j} | \theta_i, \sigma^2] = Var\left[\sum_{i'=1}^{n_j} y_{i',j} / n | \theta_i, \sigma^2\right] \quad (2)$$

$$= \frac{1}{n^2} Var\left[\sum_{i'=1}^{n_j} y_{i',j} | \theta_i, \sigma^2\right] \quad (3)$$

$$= \frac{1}{n^2} \sum_{i'=1}^{n_j} Var[y_{i',j} | \theta_i, \sigma^2] \quad (4)$$

$$= \frac{1}{n} Var[y_{i',j} | \theta_i, \sigma^2] \quad (5)$$

$$= \frac{\sigma^2}{n} \quad (6)$$

$$Cov[y_{i_1,j}, y_{i_2,j} | \theta_j, \sigma^2] = E[y_{i_1,j} y_{i_2,j}] - E[y_{i_1,j}] E[y_{i_2,j}] \quad (7)$$

$$= E[y_{i_1,j}] E[y_{i_2,j}] - E[y_{i_1,j}] E[y_{i_2,j}] \quad (8)$$

$$= 0 \quad (9)$$

² $Var(Y) = E[Var(Y|X)] + Var(E[Y|X])$

$$\text{Var}[y_{i,j}|\mu, \tau^2] = E(\text{Var}[y_{i,j}|\mu, \tau^2, \theta, \sigma^2]|\mu, \tau^2) + \text{Var}(E[y_{i,j}|\mu, \tau^2, \theta, \sigma^2]|\mu, \tau^2) \quad (10)$$

$$= E(\sigma^2|\mu, \tau^2) + \text{Var}(\theta|\mu, \tau^2) \quad (11)$$

$$= \sigma^2 + \tau^2 \quad (12)$$

$$\text{Var}[\bar{y}_{.,j}|\mu, \tau^2] = E(\text{Var}[\bar{y}_{.,j}|\mu, \tau^2, \theta, \sigma^2]|\mu, \tau^2) + \text{Var}(E[\bar{y}_{.,j}|\mu, \tau^2, \theta, \sigma^2]|\mu, \tau^2) \quad (13)$$

$$= E\left(\frac{\sigma^2}{n}|\mu, \tau^2\right) + \text{Var}(\theta|\mu, \tau^2) \quad (14)$$

$$= \frac{\sigma^2}{n} + \tau^2 \quad (15)$$

$$\text{Cov}[y_{i_1,j}, y_{i_2,j}|\mu, \tau^2] = E(\text{Cov}[y_{i_1,j}, y_{i_2,j}|\theta, \sigma^2, \mu, \tau^2]|\mu, \tau^2) \quad (16)$$

$$+ \text{Cov}(E[y_{i_1,j}|\theta, \sigma^2, \mu, \tau^2], E[y_{i_2,j}|\theta, \sigma^2, \mu, \tau^2]|\mu, \tau^2) \quad (17)$$

$$= 0 + \text{Cov}(\theta, \theta|\mu, \tau^2) \quad (18)$$

$$= E[\theta^2|\mu, \tau^2] - E[\theta|\mu, \tau^2]^2 \quad (19)$$

$$= \text{Var}(\theta|\mu, \tau^2) \quad (20)$$

$$= \tau^2 \quad (21)$$

(d)

Now assume we have a prior $p(\mu)$ for μ . Using Bayes' rule, show that

$$p(\mu|\theta_1, \dots, \theta_m, \sigma^2, \tau^2, y_1, \dots, y_m) = p(\mu|\theta_1, \dots, \theta_m, \tau^2)$$

Interpret in words what this means.

$$p(\mu|\theta_1, \dots, \theta_m, \sigma^2, \tau^2, y_1, \dots, y_m) = \frac{p(\sigma^2, y_1, \dots, y_m|\mu, \theta_1, \dots, \theta_m, \tau^2)p(\mu|\theta_1, \dots, \theta_m, \tau^2)}{p(\sigma^2, y_1, \dots, y_m|\theta_1, \dots, \theta_m, \tau^2)} \quad (22)$$

$$= p(\mu|\theta_1, \dots, \theta_m, \tau^2) \quad (23)$$

where $p(\sigma^2, y_1, \dots, y_m|\mu, \theta_1, \dots, \theta_m, \tau^2) = p(\sigma^2, y_1, \dots, y_m|\theta_1, \dots, \theta_m, \tau^2)$ because knowing μ doesn't provide more information when $\theta_1, \dots, \theta_m$ are known.

Problem 8.3

Hierarchical modeling: The files `school1.dat` through `school8.dat` give weekly hours spent on homework for students sampled from eight different schools. Obtain posterior distributions for the true means for the eight different schools using a hierarchical normal model with the following prior parameters:

$$\mu_0 = 7, \gamma_0^2 = 5, \tau_0^2 = 10, \eta_0 = 2, \sigma_0^2 = 15, \nu_0 = 2$$

```
dsch = Dict()
nsch = 8
for i in 1:nsch
    dsch[i] = readlm("data/school$i.dat")
end
```

(a)

Run a Gibbs sampling algorithm to approximate the posterior distribution of $\{\theta, \sigma^2, \mu, \tau^2\}$. Assess the convergence of the Markov chain, and find the effective sample size for $\{\sigma^2, \mu, \tau^2\}$. Run the chain long enough so that the effective sample sizes are all above 1000.

```
# Prior
0 = 7.
0^2 = 5.
0^2 = 10.
0 = 2.
0^2 = 15.
0 = 2.

# Data
ns = [ length(dsch[i]) for i in 1:nsch]
n = sum(ns)
m = length(dsch)
ys = [mean(dsch[i]) for i in 1:nsch]
s^2s = [ (ns[i] - 1)^-1 * sum( (dsch[i] .- ys[i]).^2) for i in 1:nsch]
# posterior

function ^2_pos(m, 0, v, , 0^2)
    ths^2 = sum([ ( - )^2 for in v])
    = (m + 0)* 0.5
```

```

        = (ths2 + 0*02)/2
    return InverseGamma( , )
end

function 2_pos(n, 0, 02, ns, s2s,  $\bar{y}$ s, s)
    = (n+0)/2
    =( sum((ns .- 1) .* s2s .+ ns .* ( $\bar{y}$ s .- s).^2) + 0*02)/2
    return InverseGamma( , )
end

function _pos(m, 2, s, 02, 0)
    m2 = (m/2 + 1/02)-1
    - = mean( s)

    a = m2*(m*-/2 + 0/2)
    return Normal(a, m2)
end

function _pos(2,  $\bar{y}$ , n, 2, )
    2 = (n/2 + 1/2)-1
    a = 2*(n* $\bar{y}$ /2 + /2)
    return Normal(a, 2)
end

"""
Effective Sample Size
"""
function ess(v)
    n = length(v)
    c = sum(autocov(v, collect(1:n-1)))
    return n/(1+2*c)
end

# Sampling
smp = 4000
2s = zeros(smp)
2s = zeros(smp)
s = zeros(smp)
s = zeros(smp, m)

```

```

2s[1] = rand(InverseGamma( 0/2, 0* 02/2))
2s[1] = rand(InverseGamma( 0 /2, 0 * 02/2))
s[1] = rand(Normal( 0, 02))
# s[1,:] = [rand( _pos( 2s[1],  $\bar{y}_s[i]$ , ns[i], 2s[1], s[1])) for i in 1:m]
s[1,:] = rand(Normal(s[1], 2s[1]), m)




for s in 2:smp
    2s[s] = rand( 2_pos(n, 0, 02, ns, s2s,  $\bar{y}_s$ , s[s-1,:]))
    2s[s] = rand( 2_pos(m, 0, s[s-1,:], s[s-1], 02))
    s[s,:] = [rand( _pos( 2s[s-1],  $\bar{y}_s[i]$ , ns[i], 2s[s-1], s[s-1])) for i in 1:m]
    s[s] = rand( _pos(m, 2s[s-1], s[s-1,:], 02, 0))
end

for i in [ 2s, 2s, s, s]
    plot(i)
end

p1 = plot( 2s[2:end], label=" 2 ")
p2 = plot( 2s[2:end], label=" 2 ")
p3 = plot( s[2:end], label=" ")

plot(p1, p2, p3)

```

- Effective Sample Size

- τ^2

```
ess(  $\tau^2$ s)
```

1523.525

- σ^2

```
ess(  $\sigma^2$ s)
```

1234.234

- μ

```
ess(  $\mu$ s)
```

1045.242

(b)

Compute posterior means and 95% confidence regions for $\{\sigma^2, \mu, \tau^2\}$. Also, compare the posterior densities to the prior densities, and discuss what was learned from the data.

- σ^2

```
quantile( 2s, [0.025, 0.5, 0.975])
```

3-element Vector{Float64}:

```
11.492557006363636
14.15210768607951
17.713837318481733
```

- μ

```
quantile( s, [0.025, 0.5, 0.975])
```

3-element Vector{Float64}:

```
5.3131625605052655
7.688237183428129
8.736460757511468
```

- τ

```
quantile( 2s, [0.025, 0.5, 0.975])
```

3-element Vector{Float64}:

```
1.8676028002352072
4.466482143729323
15.698240842165873
```

```
pu = density( s, label=" ")
pt = density( 2s, label=" 2 ")
ps = density( 2s, label=" 2 ")

plot!(pu, Normal( 0 , 02), label= " 0")
plot!(pt, InverseGamma( 0 /2, 0 * 02/2), label= " 2 0")
```


```

plot!(ps, InverseGamma(0/2, 0*0^2/2), label= "σ²0")

xlims!(pt, 0,30)
xlims!(ps, 0,30)

plot(pu, pt, ps)

```



Estimations of μ and τ are similar in prior and posterior. However, σ^2 is different.

(c)

Plot the posterior density of $R = \frac{\tau^2}{\sigma^2 + \tau^2}$ and compare it to a plot of the prior density of R . Describe the evidence for between-school variation.

```

²_prs = rand(InverseGamma(0/2, 0*0^2/2), 1000)
²_prs = rand(InverseGamma(0/2, 0*0^2/2), 1000)


```

```

R_prs =  $\sigma^2_{\text{prs}} / (\sigma^2_{\text{prs}} + \sigma^2_{\text{prs}})$ 
R_pos =  $\sigma^2_{\text{s}} / (\sigma^2_{\text{s}} + \sigma^2_{\text{s}})$ 

pr = density(R_prs, label="R Prior", xlabel="R", ylabel="density")
density!(pr, R_pos, label="R posterior")

```



R represents the quantity of variance in between-group. The prior is not certain about the specific quantity, but after applying posterior inference. The posterior probability of R is peaked and more certain about the value is around 0.3.

(d)

Obtain the posterior probability that θ_7 is smaller than θ_6 , as well as the posterior probability that θ_7 is smaller than of all the θ 's.

- $p(\theta_7 \text{ is smaller than } \theta_6)$

```

mean(s[:,7] .< s[:,6])

```

0.53075

- $p(\theta_7)$ is smaller than of all the θ 's)

```
res = zeros(size(s)[1])
for i in 1 : size(s)[1]
    if argmin(s[i,:]) == 7
        res[i] = 1
    end
end
mean(res)
```

0.339

(e)

Plot the sample averages $\bar{y}_1, \dots, \bar{y}_8$ against the posterior expectations of $\theta_1, \dots, \theta_8$, and describe the relationship. Also compute the sample mean of all observations and compare it to the posterior mean of μ .

```
psmp = scatter(ȳs, mean(s, dims = 1)[1,:], xlabel="Sample Average", ylabel= "Posterior Ex")
hline!(psmp, [mean(s)], label="posterior mean (n)")
hline!(psmp, [sum(ȳs .* ns)/n], label="Pooled sample mean ( )" )
```

