

Homework 9

Shao-Ting Chiu (UIN:433002162)

11/17/22

Table of contents

Computational Environment	1
Description	2
Problem 9.1	2
(a)	2
(b)	6
Problem 9.2	7
(a)	7
(b)	10

Computational Environment

```
using Pkg
Pkg.activate("hw9")
using Distributions
using DataFrames
using Plots
using DelimitedFiles
using LinearAlgebra
using Statistics
using ProtoStructs
using CSV
import Random
Random.seed!(2022)
```

```
Activating project at `~/Documents/GitHub/STAT638_Applied-Bayes-Methods/hw/hw9`
```

```
Random.TaskLocalRNG()
```

Description

- Course: STAT638, 2022 Fall
- Deadline: 2022/11/17, 12:00 pm

Read Chapter 9 in the Hoff book. Then do Problems 9.1 and 9.2 in Hoff.

For both regression models, please include an intercept term (β_0).

In 9.1(b), please replace “max” by “min”. (This is not listed in the official book errata, but appears to be a typo.)

For 9.2, the azdiabetes.dat data are described in *Exercise 6* of Chapter 7 (see errata).

- Note: This PDF file is generated by Quarto and LuaLaTeX. There is unsolved problem to display the special character in the source code. Thus, I leave the html version here for reference that displays the complete source code:

https://stchiu.quarto.pub/stat_638_hw_9/

Problem 9.1

Extrapolation: The file `swim.dat` contains data on the amount of time in seconds, it takes each of four high school swimmers to swim 50 yards. Each swimmer has 6 times, taken on a biweekly basis.

(a)

Perform the following data analysis for each swimmer separately:

1. Fit a linear regression model of swimming time as the response and week as the explanatory variable. To formulate your prior, use the information that competitive times for this age group generally range from 22 to 24 seconds.
 2. For each swimmer j , obtain a posterior predictive distribution for Y_j^* , their time if they were to swim 2 weeks from the last recorded time.
- Suppose a linear model

$$Y = X\beta + \epsilon$$

$$Y_i = x_{i,1}\beta_1 + x_{i,2}\beta_2 + \epsilon_i$$

- $Y = \begin{bmatrix} Y_1 \\ \vdots \\ Y_6 \end{bmatrix}$. A swimmer's record of 6. Series in time
- $X = \begin{bmatrix} x_{1,1} & x_{1,2} \\ \vdots & \vdots \\ x_{6,1} & x_{6,2} \end{bmatrix}$
 - $x_{j,1}$: j th record with swim score in the range of 22 to 24 second
 - $x_{j,2}$: Weeks of training
- $\beta = \begin{bmatrix} \beta_1 \\ \beta_2 \end{bmatrix}$.
 - $\mu_0 = \begin{bmatrix} 23 \\ 0 \end{bmatrix}$
 - * The prior expectation of intercept of y is 23.
 - $\beta_0 \sim N_p(\mu_0, \Sigma_0)$.
 1. FCD: $\beta|y, \sigma^2 \sim N_p(\beta_n, \Sigma_n)$
 2. $\Sigma_n^{-1} = \Sigma_0^{-1} + \frac{X^T X}{\sigma^2}$
 3. $\beta_n = \Sigma_n(\Sigma_0^{-1}\beta_0 + \frac{X^T y}{\sigma^2})$

Prior setting

- $\Sigma_0 = \begin{bmatrix} 0.1 & 0 \\ 0 & 0.1 \end{bmatrix}$
 - There is uncertainty about β estimation.
 - Covariance of time and intercept is believe as 0
- $\sigma^2 \sim IG(\nu_0/2, \nu_0\sigma_0^2/2)$
 - FCD: $\sigma^2|y, \beta \sim IG((\nu_0 + n)/2, (\nu_0\sigma_0^2) + SSR(\beta)/2)$
- $SSR(\beta) = (y - X\beta)^T(y - X\beta)$

```
ys = readdlm("data/swim.dat")
```

```
4×6 Matrix{Float64}:
23.1  23.2  22.9  22.9  22.8  22.7
23.2  23.1  23.4  23.5  23.5  23.4
22.7  22.6  22.8  22.8  22.9  22.8
23.7  23.6  23.7  23.5  23.5  23.4
```

```

"""
Problem 9.1 (a)
"""

@proto struct SwimmingModel
    S = 1000 # Number of sampling
    # Data
    y
    n = length(y) # number of records
    # Model
    X = hcat( ones(n), collect(0:2:10) )
    p = size(X)[2]
    # Prior
    = MvNormal([23., 0.], [0.1 0; 0 0.1])
    = 1.
    ^2 = 0.2
end

function SSR( , y, X)
    ssrV = (y - X*)' * (y - X*)
    return sum(ssrV)
end

function _FCD( ^2, m::SwimmingModel)
    Σ = ( m. .Σ^-1 + m.X' * m.X / ^2 )^-1
    = Σ * ( m. .Σ^-1 * m. . + m.X' * m.y / ^2 )
    return MvNormal(vec( ), Hermitian(Σ))
end

function ^2_FCD( , m::SwimmingModel)
    = (m. . + m.n)/2
    = (m. * m. ^2) + SSR( , m.y, m.X)
    return InverseGamma( , )
end

function pred(X, m::SwimmingModel)
    # Sampling vector
    smp = zeros(m.S, length(m. .))
    ^2smp = zeros(m.S)
    y = zeros(m.S)
    # Init

```

```

    smp[1,:] = rand(m. )
    ^2smp[1] = m. ^2
    y[1] = m.y[1]
    for i in 2:m.S
        smp[i,:] = rand(_FCD(^2smp[i-1], m))
        ^2smp[i] = rand(^2_FCD(smp[i-1,:], m))

        # Predict
        y[i] = smp[i,:]' * X + rand(Normal(0., ^2smp[i]))
    end

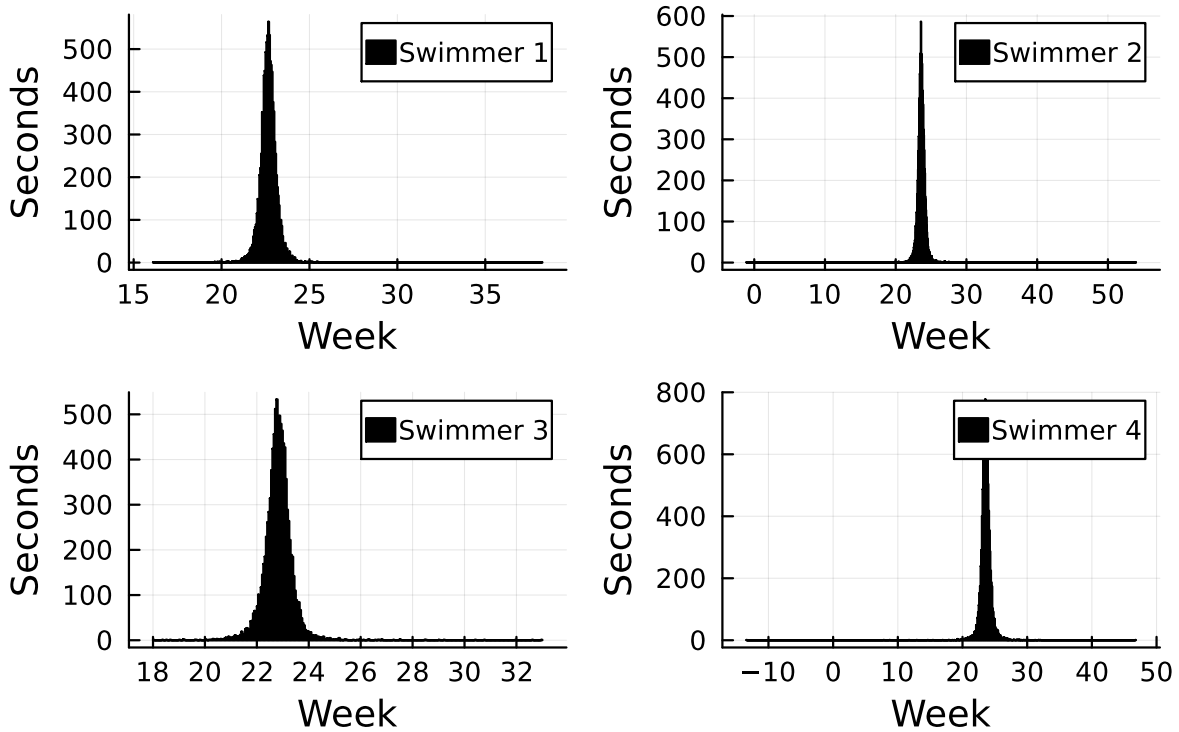
    return (y=y, smp=smp, ^2=^2smp)
end

j_swim = 1
ms = [ SwimmingModel(y = hcat(ys[i,:]), S=10000 ) for i in 1:size(ys)[1] ]
ys_pred = zeros(size(ys)[1], ms[1].S)
X_pred = [1,12]

for i in eachindex(ms)
    ys_pred[i,:] = pred([1,12], ms[i]).y
end

## Plotting
p = [histogram(ys_pred[i,:], label="Swimmer $i", color="black",
    xlabel="Week", ylabel="Seconds"
    ) for i in 1:size(ys)[1]]
plot(p...)

```



(b)

The coach of the team has to decide which of the four swimmers will compete in a swimming meet in 2 weeks. Using your predictive distributions, compute $Pr(Y_j^* = \max\{Y_1^*, \dots, Y_4^*\} | Y)$ for each swimmer j , and based on this make a recommendation to the coach.

```
am = argmax(ys_pred, dims=1)

y_count = zeros(1, size(ys)[1])

for a in am
    y_count[a[1]] += 1
end

pmax = vec(y_count ./ length(am))

## Recommendation
ds = DataFrame( Dict("Swimmer"=> collect(1:size(ys)[1]), "Pr(Y_i is max)" => pmax ))
```

	Pr(Y_i is max)	Swimmer
	Float64	Int64
1	0.023	1
2	0.4773	2
3	0.0424	3
4	0.4573	4

Swimmer 2 is the most probable winner.

Problem 9.2

Model selection: As described in Example 6 of Chapter 7, the file `azdiabetes.dat` contains data on health-related variables of a population of 532 women. In this exercise we will be modeling the conditional distribution of glucose level (`glu`) as a linear combination of the other variables, excluding the variable `diabetes`.

(a)

Fit a regression model using the g -prior with $g = n$, $\nu_0 = 2$ and $\sigma_0^2 = 1$. Obtain posterior confidence intervals for all of the parameters.

```
data = readlm("data/azdiabetes.dat")
```

533×8 Matrix{Any}:

	"npreg"	"glu"	"bp"	"skin"	"bmi"	"ped"	"age"	"diabetes"
5	86	68	28	30.2	0.364	24	"No"	
7	195	70	33	25.1	0.163	55	"Yes"	
5	77	82	41	35.8	0.156	35	"No"	
0	165	76	43	47.9	0.259	26	"No"	
0	107	60	25	26.4	0.133	23	"No"	
5	97	76	27	35.6	0.378	52	"Yes"	
3	83	58	31	34.3	0.336	25	"No"	
1	193	50	16	25.9	0.655	24	"No"	
3	142	80	15	32.4	0.2	63	"No"	
2	128	78	37	43.3	1.224	31	"Yes"	
0	137	40	35	43.1	2.288	33	"Yes"	
9	154	78	30	30.9	0.164	45	"No"	
12	100	84	33	30.0	0.488	46	"No"	
1	147	94	41	49.3	0.358	27	"Yes"	

3	187	70	22	36.4	0.408	36	"Yes"
1	121	78	39	39.0	0.261	28	"No"
3	108	62	24	26.0	0.223	25	"No"
0	181	88	44	43.3	0.222	26	"Yes"
1	128	88	39	36.5	1.057	37	"Yes"
2	88	58	26	28.4	0.766	22	"No"
9	170	74	31	44.0	0.403	43	"Yes"
10	101	76	48	32.9	0.171	63	"No"
5	121	72	23	26.2	0.245	30	"No"
1	93	70	31	30.4	0.315	23	"No"

```

dt = data[1:end , 1:end-1]
y = float.(dt[2:end, 2])
X = float.(dt[2:end, 1:end .!= 2])
ns = data[1,1:end-1]
ns = ns[1:end .!=2]

@proto struct DiabetesModel
    S = 500 # Number of sampling
    # Data
    y
    X
    n = length(y) # number of records
    p = size(X)[2]
    # Model
    # Prior
    g = n # g prior
    = 2.
    ^2 = 1.
end

function _FCD( ^2, m::DiabetesModel)
    return _FCD( ^2, m.g, m.X, m.y)
end

function _FCD( ^2, g, X, y)
    Σ = g/(g+1) * ^2 * (X'X)^-1
    = g/(g+1) * ^2 * (X'X)^-1
    = MvNormal( , Hermitian(Σ))
    return
end

```



```

function beta_hat(beta_hat, y, X)
    return beta_hat * (X'X)^-1 * (X'y / beta_hat)
end

function beta_hat_FCD(m::DiabetesModel)
    return beta_hat_FCD(m.beta_hat, m.n, m.X, m.y, m.g)
end

function beta_hat_FCD(beta_hat, n, X, y, g)
    beta_hat = beta_hat + n / 2.
    beta_hat = (beta_hat * beta_hat + SSR(X, y, g)) / 2.
    beta_hat = InverseGamma(beta_hat, )
    return beta_hat
end

function SSR(m::DiabetesModel)
    return SSR(m.X, m.y, m.g)
end

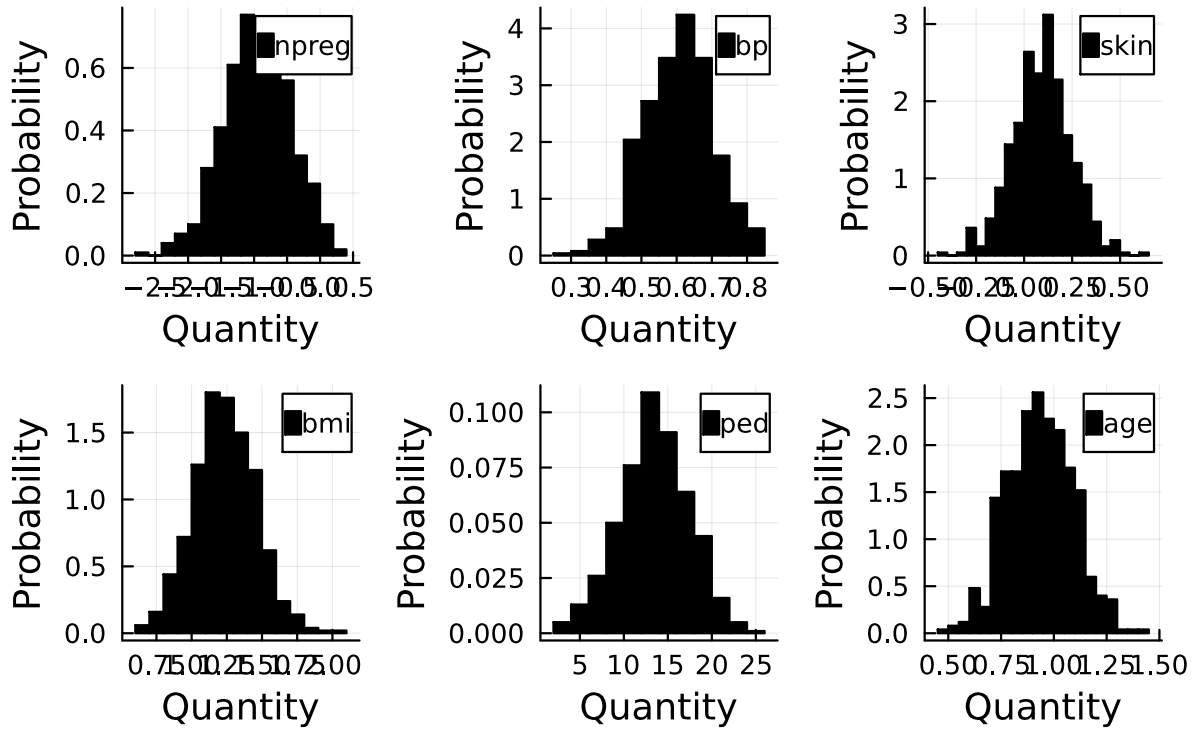
function SSR(X, y, g)
    return y'(I - g/(g+1)*X*(X'X)^-1*X')*y
end

m = DiabetesModel(y=y, X=X)
beta_hat_smp = zeros(m.S, 1)
smp = zeros(m.S, size(m.X)[2])

for i in 1:m.S
    beta_hat_smp[i] = rand(beta_hat_FCD(m))
    smp[i,:] = rand(beta_hat_FCD(beta_hat_smp[i], m))
end

ps = [histogram(smp[:,i], xlabel="Quantity",
    ylabel="Probability", normalize=true, label="$(ns[i])", color="black")
    for i in 1:m.p]
plot(ps...)

```



(b)

Perform the model selection and averaging procedure described in Section 9.3. Obtain $Pr(\beta_j \neq 0|y)$, as well as posterior confidence intervals for all of the parameters. Compare to the results in part (a).

$$\{y|X_{z(k)}, \beta_{z(k)}, \sigma^2\} \sim \text{multivariate normal } (X_{z(k)}\beta_{z(k)}, \sigma^2 I)$$

```
m = DiabetesModel(y=y, X=X, S=1000)

function ²_FCD(m::DiabetesModel, zs)
    Xz = @view m.X[1:end, Bool.(zs)]
    return ²_FCD(m. , m. ², m.n, Xz, m.y, m.g)
end

function _FCD(², m::DiabetesModel, zs)
    Xz = @view m.X[1:end, Bool.(zs)]
```

```

        return _FCD(z^2, m.g, Xz, m.y)
    end

function y_margin(z^2, m::DiabetesModel, zs)
    = m.
    n = m.n
    y = m.y
    g = m.g
    pz = sum(zs)
    Xz = @view m.X[1:end, Bool.(zs)]
    ssr = SSR(Xz, y, g)

    pyl = (pz/2.)log(1. +g) + ( /2.)*log(z^2) - (( +n)/2)log(( * z^2 + ssr))
    return pyl
end

function z_FCD(i , z^2, zsmpl, nSmp,m::DiabetesModel)
    zs = zsmpl[nSmp,:]
    pj1 = sum(zsmpl[1:nSmp, i]) / length(zsmpl[1:nSmp, i])
    pj0 = 1. - pj1
    pj1_FCD_l = pj1 * y_margin(z^2, m, ones(length(zs)))
    pj0_FCD_l = pj0 * y_margin(z^2, m, zs)
    O = exp(pj0_FCD_l - pj1_FCD_l)
    return Bernoulli( 1/(1+O))
end

zsmpl = ones(m.S, size(m.X)[2])
^2smp = zeros(m.S, 1)
smp = zeros(m.S, size(m.X)[2])

^2smp[1] = 0.1
# Gibbs sampling
for i in 2:m.S
    for j in Random.shuffle(1:m.p)
        zsmpl[i, j] = rand(z_FCD(j, ^2smp[i-1], zsmpl, i-1, m))
    end

    ^2smp[i] = rand(^2_FCD(m, zsmpl[i,:]))
    smp[i, Bool.(zsmpl[i,:])] = rand(_FCD(^2smp[i], m, zsmpl[i,:]))
end
end

```

```

sum(zsmp, dims=1)/size(zsmp)[1]
prB = 1. - vec(sum(zsmp, dims=1))./size(zsmp)[1]
DataFrame(Dict( "Bi"=> ns, "Pr(Bi != 0 |y)"=> prB ))


```

	Bi	Pr(Bi != 0 y)
	Any	Float64
1	npreg	0.499
2	bp	0.499
3	skin	0.499
4	bmi	0.499
5	ped	0.499
6	age	0.5

```

inds = prB .>= 0.5
b_select = prB[inds]
ps = [histogram( smp[:, i], xlabel="Quantity",
                ylabel="Probability",normalize=true, label="$(ns[i])", color="black")
      for i in findall(inds .== 1)]
plot(ps...)

```



```
DataFrame(Dict("Parameters"=> ns[inds],
               "Confidence interval"=> [quantile(smp[:,i], [0.25, 0.975]) for i in finda
```

	Confidence interval	Parameters
	Array...	Any
1	[0.0, 1.1847]	age

Conclusion

There might be some bugs in FCD formulation. The current results show that the **age** is the only factor remained after model selection. The distribution is far different from what it is in part (a).